Lecture 6

Brad McNeney

2017-02-07

Databases

Relational database

- A relational database is a collection of tables.
    - As statisticians, we think of rows (a.k.a. records) as sampled units, and columns as the variables measured on the units.
- Rows of one table are related to rows of others by keys
    - Each table has a key (primary key) that uniquely identifies the rows.
    - Including the key of another table (a foreign key) allows us to link (relate) records of the two tables.

Simple example database

- ▶ STATION table

ID	City	State	Lat_N	Long_W
13	Phoenix	AZ	33	112
44	Denver	CO	40	105
66	Caribou	ME	47	68

- ▶ STATS table (primary key suppressed; ID is foreign key)

ID	Month	Temp_F	Rain_I
13	1	57.4	0.31
13	7	91.7	5.15
44	1	27.3	0.18
44	7	74.8	2.11
66	1	6.7	2.1
66	7	65.8	4.52

Database software terminology

- Software refered to as relational database management systems (RDBMS).
    - Example implementations: MySQL, SQLite
- The language for querying the database is structured query language (SQL).
    - Similar implementations in all major RDBMS
- The database resides on a server, which we access with a client.
    - Large databases will be on special-purpose remote servers.
    - Client runs on your computer.
    - However, client/server model usually involves layers of security that make access difficult.
    - We'll work mostly with SQLite – server runs on your computer.

Why databases in R?

- The data is stored in a RDBMS.
  - Our focus today
- Work with datasets that are too large to be stored in memory
  - Discuss later

SQLite databases in R

- ▶ The `RSQLite` package contains an SQLite client *and* server we can use from R

  - ▶ SQLite is an open-source RDBMS engine bundled with RSQLite
  - ▶ Interface with DB engine (connect, write to, etc.) *via* a common interface called DBI
  - ▶ Install `RSQLite` and `DBI` before starting
  - ▶ Load DBI; access functions from `RSQLite` with the `::` operator.

```r
library(DBI)
mydb <- dbConnect(RSQLite::SQLite(), "my-db.sqlite")
dbDisconnect(mydb)
```

```
## [1] TRUE
```

```
#> [1] TRUE
unlink("my-db.sqlite")
```

Notes

- The first line, `mydb <- dbConnect(RSQLite::SQLite(), "my-db.sqlite")`, creates a "connection" to an SQLite database stored in the file "my-db.sqlite"
  - `SQLite()` is called the "driver". It handles all the RDBMS-specific details of how the client and server communicate.
  - The following argument is the file name for the database. Other drivers need details like username and password to connect.
  - Initially the database is empty
- The second line, `dbDisconnect(mydb)`, disconnects from the database, but does not remove it.
- Remove with `unlink()` (a base R command).
- `dbDisconnect()` and all other commands we will use of the form `db*` are generics from DBI

Aside: MySQL database driver in R

- Wasn't able to get easy access to a MySQL database on campus.
- In case it is of later use, here are the arguments to `dbConnec()` for connecting to a password-protected database for Stat 341.

```
con <- dbConnect(MySQL(), user='stataccess',
                 password='<password goes here>',
                 dbname='stat341',
                 host='muncho.its.sfu.ca')
dbListTables(con)
```

Creating `STATION` and `STATS` tables

```
wdb <- dbConnect(RSQLite::SQLite(), "wdb.sqlite")
STATION <- data.frame(ID=c(13,44,66),
    City = c("Phoenix","Denver","Caribou"),
    State = c("AZ","CO","ME"),
    Lat_N = c(33,40,47),
    Long_W = c(112,105,68))
STATS <- data.frame(row = 1:6,
    ID = c(13,13,44,44,66,66),
    Month = c(1,7,1,7,1,7),
    Temp_F = c(57.4,91.7,27.3,74.8,6.7,65.8),
    Rain_I = c(0.31,5.15,0.18,2.11,2.1,4.52))
dbWriteTable(wdb,name='STATION', value = STATION, overwrite=TRUE)
```

```
## [1] TRUE
```

```
dbWriteTable(wdb,name='STATS', value = STATS, overwrite=TRUE)
```

```
## [1] TRUE
```

```
dbListTables(wdb)
```

```
## [1] "STATION" "STATS"
```

Notes

- The `overwrite=TRUE` argument to the `dbWriteTable()` commands is necessary for this demo, to allow re-knitting the document, but is not necessary in general.
- Now that we have an example database, we can see how to use SQL to extract data.
- Will interleave SQL tutorial and `DBI`.

SQL: Retrieve data from a single table

- The SQL `SELECT` statements:
  - `SELECT * FROM STATION;`
  - `SELECT City, Lat_N FROM STATION;`
  - `SELECT Month, Temp_F, Rain_I FROM STATS;`
- Filtering with the `WHERE` clause:
  - `SELECT * from STATION WHERE Lat_N >= 40;`
- Trailing `;` required for most SQL clients, but not R DBI.

DBI: Example data retrieval.

```
dbGetQuery(wdb,"SELECT * from STATION")
```

```
##   ID    City State Lat_N Long_W
## 1 13 Phoenix    AZ    33    112
## 2 44  Denver    CO    40    105
## 3 66 Caribou    ME    47     68
```

```
dbGetQuery(wdb,"SELECT City, Lat_N from STATION")
```

```
##      City Lat_N
## 1 Phoenix    33
## 2  Denver    40
## 3 Caribou    47
```

```
dbGetQuery(wdb,"SELECT * from STATION WHERE Lat_N>=40")
```

```
##   ID    City State Lat_N Long_W
## 1 44  Denver    CO    40    105
## 2 66 Caribou    ME    47     68
```

DBI: Notes on `dbGetQuery()`

- `dbGetQuery()` calls three functions:
  1. `dbSendQuery()` sends the query to the DB,
  2. `dbFetch()` fetches the "result set", and
  3. `dbClearResult()` frees memory and other resources associated with the result set.

- If the result set is too large to fit in memory, you can split the fetching into batches.

DBI: Batched queries

```r
rs <- dbSendQuery(wdb,"SELECT * FROM STATS")
while (!dbHasCompleted(rs)) {
  df <- dbFetch(rs, n = 2) # use n to set size of subset
  print(df)
}
```

```
##   row ID Month Temp_F Rain_I
## 1   1 13     1   57.4   0.31
## 2   2 13     7   91.7   5.15
##   row ID Month Temp_F Rain_I
## 1   3 44     1   27.3   0.18
## 2   4 44     7   74.8   2.11
##   row ID Month Temp_F Rain_I
## 1   5 66     1    6.7   2.10
## 2   6 66     7   65.8   4.52
```

```r
dbClearResult(rs)
```

```
## [1] TRUE
```

DBI: Parametrized queries

- Can pass the same query with several different values of a parameter `x`.
  - Bind a value to the parameter with dbBind()

```
rs <- dbSendQuery(wdb,"SELECT * FROM STATION WHERE Lat_N >= :x")
dbBind(rs,param = list(x=40))
dbFetch(rs)
```

```
##   ID    City State Lat_N Long_W
## 1 44  Denver    CO    40    105
## 2 66 Caribou    ME    47     68
```

```
dbBind(rs,param=list(x=45))
dbFetch(rs)
```

```
##   ID    City State Lat_N Long_W
## 1 66 Caribou    ME    47     68
```

```
dbClearResult(rs)
```

```
## [1] TRUE
```

SQL: Joining tables

- The purpose of related tables is to reduce redundancy.

  - For example, all the info on the stations appears once in the STATION table, and need not be repeated in the STATS table.

- But what if we need the info on the stations and the weather data? Need to "join" tables.

- Simplest join (inner join): SELECT * from STATION, STATS WHERE STATION.ID=STATS.ID

DBI: Example joins

```
dbGetQuery(wdb,"SELECT * from STATION, STATS WHERE STATION.ID=STATS.ID")
```

```
##   ID    City State Lat_N Long_W row ID Month Temp_F Rain_I
## 1 13 Phoenix    AZ    33    112   1 13     1   57.4   0.31
## 2 13 Phoenix    AZ    33    112   2 13     7   91.7   5.15
## 3 44  Denver    CO    40    105   3 44     1   27.3   0.18
## 4 44  Denver    CO    40    105   4 44     7   74.8   2.11
## 5 66 Caribou    ME    47     68   5 66     1    6.7   2.10
## 6 66 Caribou    ME    47     68   6 66     7   65.8   4.52
```

```
queryp1 <- "SELECT City, State, Month, Rain_I from STATION, STATS"
queryp2 <- "WHERE STATION.ID=STATS.ID AND Month = 1 AND Lat_N >= 40"
dbGetQuery(wdb,paste(queryp1,queryp2))
```

```
##      City State Month Rain_I
## 1  Denver    CO     1   0.18
## 2 Caribou    ME     1   2.10
```

SQL: Left joins

- The inner join returns data for cities in **both** the STATION and STATS tables.
- If we want all stations, use a left join.
    - First add another station with no data in STATS
    - Miami, FL is at Lat 26 and Long 80.
    - Give Miami station ID 77.
    - In SQL we'd add Miami and do the left join with

      ```
      INSERT INTO STATION VALUE (77,'Miami','FL',26,80)
      SELECT * FROM STATION
      LEFT JOIN STATS ON STATION.ID = STATS.ID
      ```

- Many other types of SQL joins. See [https://www.tutorialspoint.com/sqlite/sqlite_using_joins.htm] for a summary of joins in SQLite.

DBI: Adding to a table and left join

```
miami <- data.frame(ID=77,City="Miami",State="FL",Lat_N=26,Long_W=80)
dbWriteTable(wdb,name='STATION', value = miami, append=TRUE)
```

```
## [1] TRUE
```

```
qq<-"SELECT * FROM STATION LEFT JOIN STATS ON STATION.ID = STATS.ID"
qq<-"SELECT City, State, Temp_F FROM STATION LEFT JOIN STATS ON STATION
dbGetQuery(wdb,qq)
```

```
##       City State Temp_F
## 1  Denver    CO   27.3
## 2  Denver    CO   74.8
## 3 Caribou    ME    6.7
## 4 Caribou    ME   65.8
```

- ▶ Note: `append=TRUE` adds to the current table

SQL: Table indices

- A query like SELECT * FROM STATION WHERE Lat_N >= 40 requires that the RDBMS read the Lat_N value in every row of STATION and return the rows where Lat_N is 40 or more
- Such a query can be made much faster by creating an "index" on Lat_N.
    - An index is a table in the database, sorted on the indexed variable.
    - See [http://www.sqlite.org/queryplanner.html] for a nice description of how indexing columns speeds up searches.

DBI: Create an index with dbExecute()

▶ Use dbExecute() to execute queries that do not return tabular data.

```
dbExecute(wdb,"CREATE INDEX indx ON STATION(Lat_N)")
```

```
## [1] 1
```

Exercises

1. Select the stations from the state of Colorado
2. Add Vancouver as a new station in the STATION table (Use BC as the "State" and the latitude 49 and longitude 123)
3. Do an inner join that returns a table with city, state, and temperatures for July from cities at north latitude 40 or less.
4. Repeat (3) as a left join.

More Exercises

5. Do a left join that returns a table with city, state, and temperatures for July from cities at north latitude 40 or more. (Vancouver excluded because we filter on month, and month is missing for Vancouver.)

6. Do a left join that returns a table with city, state, and temperatures from cities at north latitude 40 or less.

```
vancouver <- data.frame(ID=88,City="Vancouver",State="BC",Lat_N=49,Long_W=123)
dbWriteTable(wdb,name='STATION', value = vancouver, append=TRUE)
```

```
## [1] TRUE
```

```
qq<-"SELECT * FROM STATION LEFT JOIN STATS ON STATION.ID = STATS.ID"
qq<-"SELECT City, State, Temp_F FROM STATION LEFT JOIN STATS ON STATION.ID = ST
dbGetQuery(wdb,qq)
```

```
##           City State Temp_F
## 1     Denver    CO    27.3
## 2     Denver    CO    74.8
## 3    Caribou    ME     6.7
## 4    Caribou    ME    65.8
## 5  Vancouver    BC      NA
```

Clean up

```
dbDisconnect(wdb)
```

```
## [1] TRUE
```

```
unlink("wdb.sqlite")
```

Merging, selecting and filtering on data frames

Inner join data frames with `merge()`

▶ The `merge()` function in R does SQL-like joins on data frames.

```
STATION <- rbind(STATION,miami)
merge(STATION,STATS,by="ID")
```

```
##    ID    City State Lat_N Long_W row Month Temp_F Rain_I
## 1 13 Phoenix    AZ    33    112   1     1   57.4   0.31
## 2 13 Phoenix    AZ    33    112   2     7   91.7   5.15
## 3 44  Denver    CO    40    105   3     1   27.3   0.18
## 4 44  Denver    CO    40    105   4     7   74.8   2.11
## 5 66 Caribou    ME    47     68   5     1    6.7   2.10
## 6 66 Caribou    ME    47     68   6     7   65.8   4.52
```

Left joining data frames with merge()

```
merge(STATION,STATS,by="ID",all.x=TRUE)
```

```
##    ID    City State Lat_N Long_W row Month Temp_F Rain_I
## 1 13 Phoenix    AZ    33    112   1     1   57.4   0.31
## 2 13 Phoenix    AZ    33    112   2     7   91.7   5.15
## 3 44  Denver    CO    40    105   3     1   27.3   0.18
## 4 44  Denver    CO    40    105   4     7   74.8   2.11
## 5 66 Caribou    ME    47     68   5     1    6.7   2.10
## 6 66 Caribou    ME    47     68   6     7   65.8   4.52
## 7 77   Miami    FL    26     80  NA    NA     NA     NA
```

Join functions in `dplyr`

- In `dplyr` the functions for joining are more explicitly named

```
library(dplyr)
inner_join(STATION,STATS,by="ID")
```

```
##    ID    City State Lat_N Long_W row Month Temp_F Rain_I
## 1 13 Phoenix    AZ    33    112   1     1   57.4   0.31
## 2 13 Phoenix    AZ    33    112   2     7   91.7   5.15
## 3 44  Denver    CO    40    105   3     1   27.3   0.18
## 4 44  Denver    CO    40    105   4     7   74.8   2.11
## 5 66 Caribou    ME    47     68   5     1    6.7   2.10
## 6 66 Caribou    ME    47     68   6     7   65.8   4.52
```

Left join function in `dplyr`

```
left_join(STATION,STATS,by="ID")
```

```
## ID    City State Lat_N Long_W row Month Temp_F Rain_I
## 1 13 Phoenix    AZ    33    112   1     1   57.4   0.31
## 2 13 Phoenix    AZ    33    112   2     7   91.7   5.15
## 3 44  Denver    CO    40    105   3     1   27.3   0.18
## 4 44  Denver    CO    40    105   4     7   74.8   2.11
## 5 66 Caribou    ME    47     68   5     1    6.7   2.10
## 6 66 Caribou    ME    47     68   6     7   65.8   4.52
## 7 77   Miami    FL    26     80  NA    NA     NA     NA
```

select() to select columns

- ▶ select() from dplyr can be used to select columns.

  - ▶ Can use different "helper" functions to select variables
    (help(select_helpers))

```
select(STATION,City,State)
```

```
##        City State
## 1 Phoenix    AZ
## 2  Denver    CO
## 3 Caribou    ME
## 4   Miami    FL
```

```
select(STATION,matches("L."))
```

```
##    Lat_N Long_W
## 1    33     112
## 2    40     105
## 3    47      68
## 4    26      80
```

Using `filter()` like WHERE

```r
select(STATION,matches("L.")) %>% filter(Lat_N>=40)
```

```
##   Lat_N Long_W
## 1    40    105
## 2    47     68
```

Combining join/select/filter with %>%

```
inner_join(STATION,STATS,by="ID") %>%
  select(matches("._.")) %>% filter(Lat_N>=40)
```

```
##    Lat_N Long_W Temp_F Rain_I
## 1     40    105   27.3   0.18
## 2     40    105   74.8   2.11
## 3     47     68    6.7   2.10
## 4     47     68   65.8   4.52
```

Combining join/select/filter with %>%

- After SQL it is natural to chain together joining, selecting and filtering.

```
inner_join(STATION,STATS,by="ID") %>%
  select()
```

```
## data frame with 0 columns and 6 rows
```

Exercises with `dplyr`

1. Select the stations from the state of Colorado (`filter()`)
2. Select the City and State from the stations using the `select()` command from dplyr: `select(STATION,City,State)`
3. Add Vancouver as a new station in the `STATION` data frame. Use BC as the "State" and the latitude 49 and longitude 123. (`rbind()` or `bind_rows()`)
4. Do an inner join that returns a table with city, state, and temperatures for July from cities at north latitude 40 or less. (`inner_join()` and `select()`)
5. Repeat (3) as a left join. (`left_join()` and `select()`)