

## Lecture 9

Brad McNeney

2017-03-14

## Load packages

```
library(broom)
library(ggplot2)
library(dplyr)
library(gapminder)
data(gapminder)
```

# Topics

- ▶ Exploratory data analysis, with emphasis on ggplot2 graphics, using the gapminder data.
  - ▶ Suppose we want to use information on continent, year, pop and gdpPercap to predict lifeExp.
- ▶ Base R graphics vs grid graphics
- ▶ Introduction to ggplot2

## Exploratory analyses

# Exploratory analyses

- ▶ Univariate summaries, such as means/medians, sds/IQRs, histograms, to examine distributions and identify possible measurement errors.
- ▶ Pair-wise correlations, to look for relationships between variables
- ▶ Pair-wise regression relationships and added-variable-plots
  - ▶ Trends over time deserve special attention

# Univariate Summaries

- ▶ Different summaries are appropriate for categorical and quantitative variables
  - ▶ Tabulate categorical variables
  - ▶ Five number summary for quantitative variables

```
summary(gapminder)
```

```
##           country      continent      year      lifeExp
## Afghanistan: 12 Africa :624   Min.    :1952   Min.    :23.60
## Albania      : 12 Americas:300   1st Qu.:1966   1st Qu.:48.20
## Algeria      : 12 Asia    :396   Median  :1980   Median  :60.71
## Angola       : 12 Europe  :360   Mean    :1980   Mean    :59.47
## Argentina    : 12 Oceania : 24   3rd Qu.:1993   3rd Qu.:70.85
## Australia    : 12                Max.    :2007   Max.    :82.60
## (Other)      :1632
##           pop      gdpPercap
## Min.    :6.001e+04   Min.    : 241.2
## 1st Qu.:2.794e+06   1st Qu.: 1202.1
## Median :7.024e+06   Median : 3531.8
## Mean    :2.960e+07   Mean    : 7215.3
## 3rd Qu.:1.959e+07   3rd Qu.: 9325.5
## Max.    :1.319e+09   Max.    :113523.1
##
```

## Comments on summaries

- ▶ Observations in pop and gdpPercap differ by orders of magnitude
  - ▶ May be more informative to consider transformations of these variables.
  - ▶ For example, a log-10 transformation: one-unit differences correspond to 10-fold increases.
- ▶ Aside: Which country has per-capita GDP of \$113,523? Or more generally, which observations are in, say, the top 0.1%?

```
filter(gapminder, gdpPercap > quantile(gdpPercap, 0.999))
```

```
## # A tibble: 2 × 6
##   country continent  year lifeExp    pop gdpPercap
##   <fctr>    <fctr> <int>   <dbl> <int>    <dbl>
## 1 Kuwait      Asia  1957  58.033 212846 113523.1
## 2 Kuwait      Asia  1972  67.712 841934 109347.9
```

# Univariate summaries by grouping variable

- May be of interest to do summaries of some variables stratified by a grouping variable.

```
oldops <- options(tibble.width=Inf, tibble.print_max=Inf)
gm_byContinent <- group_by(gapminder, continent)
summarize(gm_byContinent, min(lifeExp), median(lifeExp), IQR(lifeExp), mean(lifeExp))
```

```
## # A tibble: 5 × 7
##   continent `min(lifeExp)` `median(lifeExp)` `IQR(lifeExp)`
##   <fctr>      <dbl>          <dbl>          <dbl>
## 1 Africa      23.599          47.7920         12.0390
## 2 Americas    37.579          67.0480         13.2895
## 3 Asia        28.801          61.7915         18.0790
## 4 Europe      43.585          72.2410          5.8805
## 5 Oceania     69.120          73.6650          6.3475
##   `mean(lifeExp)` `sd(lifeExp)` `max(lifeExp)`
##   <dbl>          <dbl>          <dbl>
## 1 48.86533        9.150210       76.442
## 2 64.65874        9.345088       80.653
## 3 60.06490       11.864532       82.603
## 4 71.90369        5.433178       81.757
## 5 74.32621        3.795611       81.235
```


```
options(oldops)
```



# Boxplots


- Graphical representation of the five number summary for grouped data

```
ggplot(gapminder, aes(x=continent, y=lifeExp)) + geom_boxplot()
```



## Boxplots, cont.


```
ggplot(gapminder, aes(x=continent, y=gdpPercap)) + geom_boxplot()
```



- Distribution of log-transformed data may be more informative.

## Boxplots, cont.

```
ggplot(gapminder, aes(x=continent, y=gdpPercap)) +  
  coord_trans(y="log10") + geom_boxplot()
```



# Adding transformed variables to a dataset


- ▶ Above suggests we add log of gdpPercap to the dataset.
- ▶ A similar exploration of the pop variable suggests we include log of pop too.
- ▶ Will use log-base-10.

```
gapminder <- mutate(gapminder,  
  log10Pop = log10(pop),  
  log10GdpPercap = log10(gdpPercap))
```

# Histograms

- ▶ Shows the shape of distributions and can suggest possible outliers
- ▶ Stacked histograms:


```
ggplot(gapminder, aes(x=lifeExp, color=continent)) + geom_histogram()
```



# Histograms, continued

- Histograms in different plot panels, or “facets”:


```
ggplot(gapminder, aes(x=lifeExp, color=continent)) +  
  geom_histogram() + facet_grid(continent ~ .)
```



# Histograms of the explanatory variables

- May also be of interest

```
ggplot(gapminder, aes(x=log10GdpPercap, color=continent)) +  
  geom_histogram() + facet_grid(continent ~ .)
```



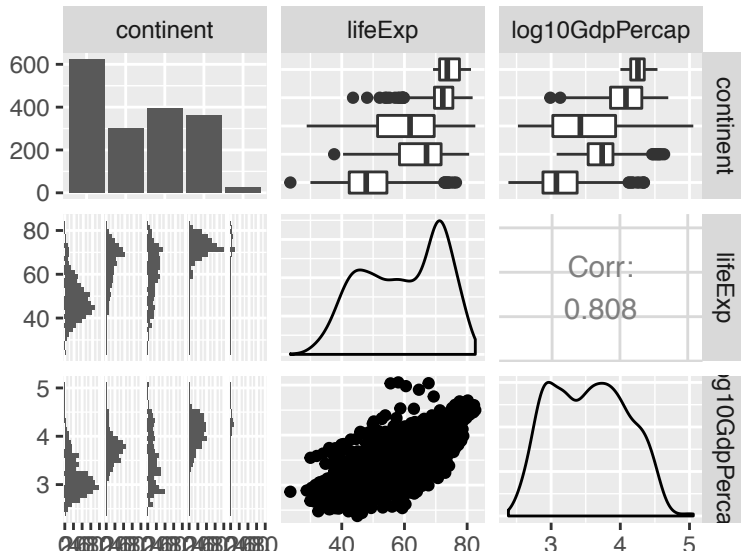
# Pairwise Regression relationships

- ▶ Though pairwise relationships don't tell the whole story, they are a useful starting point.
- ▶ The `GGally` package provides the function `ggpairs()` to facilitate this.
  - ▶ Can do all possible pairs of variables, but I find this too hard to read for more than three variables.



# Pairwise plots


```
library(GGally)
gm_sub <- select(gapminder, continent, lifeExp, log10GdpPercap)
ggpairs(gm_sub) # Cut and paste into console to see better
```



## Time trends

- ▶ Can represent time series by lines.
- ▶ There are many time series in these data – need to make each line slightly transparent to account for overplotting


```
ggplot(gapminder, aes(x=year, y=lifeExp, group=country)) +  
  geom_line(alpha=0.2)
```



## Time trends, cont.

- ▶ Can add a statistical summary, like medians at each time, or a smoother.
- ▶ Can also add colours for different continents.

```
ggplot(gapminder, aes(x=year, y=lifeExp, group=country,  
                      color=continent)) + geom_line(alpha=0.2) +  
geom_smooth(aes(x=year, y=lifeExp, group=continent), se=FALSE)
```



## Base R graphics

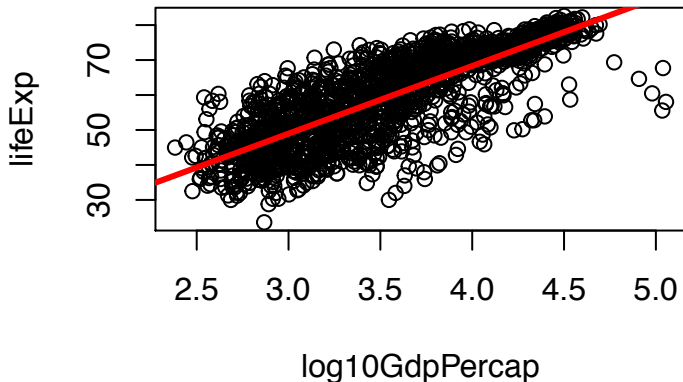
# Base R graphics

- ▶ Very serviceable graphics system capable of producing publication-quality graphs.
- ▶ Create graphics by calling functions that either produce complete plots or add to plots
- ▶ Like adding paint to a canvas

## Base R examples


```
with(gapminder,plot(log10GdpPercap,lifeExp)) # or plot(lifeExp ~ log10GdpPercap,  
title(main="life expectance vs log10 GDP percapita")  
abline(lm(lifeExp ~ log10GdpPercap,data=gapminder),col="red",lwd=3)
```

### life expectance vs log10 GDP percapita



## Base R examples


```
with(gapminder,boxplot(split(lifeExp,continent)))
```



## Base R examples

```
with(gapminder,plot(density(lifeExp)))
```

**density.default(x = lifeExp)**



N = 1704   Bandwidth = 2.625



## Base R graphics: Where to learn more

- ▶ Paul Murrell's book:  
[<https://www.stat.auckland.ac.nz/~paul/RG2e/>]
- ▶ Ross Ihaka's lectures: [<https://www.stat.auckland.ac.nz/~ihaka/787/lectures-r-graphics.pdf>]

grid graphics

# grid graphics

- ▶ grid graphics is a low-level graphics system that allows fine control of graphics elements
- ▶ Users can create multiple graphics regions, or “viewports”, that are arranged on the graphics device or nested within one another.
- ▶ Graphical objects, or “grobs” can be created and drawn on these viewports (e.g., lines, shapes).
- ▶ Grobs can be edited (e.g., change color of lines) and re-drawn

## grid graphics: Where to learn more

- ▶ Paul Murrell's book:  
[<https://www.stat.auckland.ac.nz/~paul/RG2e/>]

ggplot2 basics

# ggplot2

- ▶ ggplot2 is implemented in grid graphics
- ▶ The g's stand for Grammar of Graphics.
  - ▶ Like English grammar is the way in which words are put together to form sentences, a grammar of graphics is a way to put together basic graphical elements to make a graph.
- ▶ To understand the grammar we need to define the basic elements.
  - ▶ Start with definitions (in bold), some of which are too abstract to be useful until we get into details.
- ▶ ggplots can be built in layers, comprised of **data** a **mapping**, a **geom** and optionally **stats**
- ▶ The layers are arranged and labelled on the graph by **scales** and **coords**.
- ▶ The data can also be broken into subsets and displayed in separate graphs by a **facet** specification.

# Components of a ggplot: data and mappings

- ▶ We start with the **data** we want to visualize and a **mapping**, or aesthetic, that describes how these data map to attributes on the plot.

```
p <- ggplot(gapminder, aes(x=log10GdpPercap, y=lifeExp, color=continent))
```

- ▶ From the dataset `gapminder`, the variable `log10GdpPercap` will be mapped to y-coordinates, `lifeExp` will be mapped to the x-coordinates, and `continent` will be perceived as colours.

# Components of a ggplot: geometric objects (geoms)

- ▶ Geometric objects or **geoms** are things like points and lines that we see on the plot.

```
p2 <- p + geom_point(alpha=0.5)
```

- ▶ alpha is the transparency aesthetic, between 0 and 1, best applied directly to the geom it is to apply to

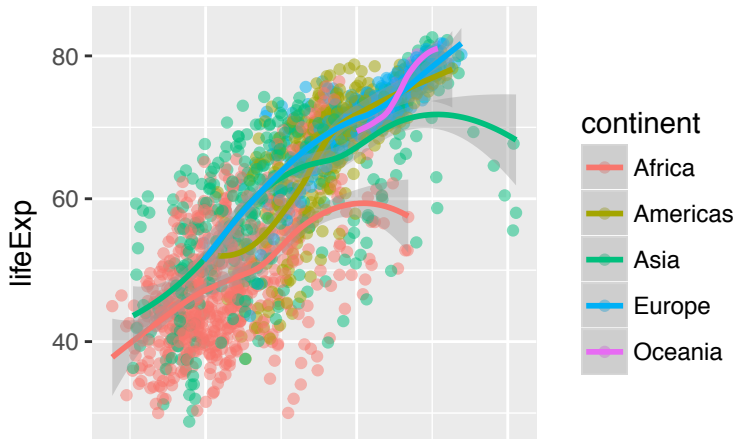


# Components of a ggplot: statistical transformations

- ▶ Statistical transformations or **stats** summarize the data; e.g., a scatterplot smoother

```
p2 + geom_smooth()
```


```
## `geom_smooth()` using method = 'loess'
```



# Components of a ggplot: scales

- ▶ The **scales** are mappings from the data to the graphics device
  - ▶ domain of `continent` is the five continents, range is the hexadecimal of the five colors represented on the graph
  - ▶ domain of `lifeExp` is 23.599 to 82.603, range is  $[0,1]$ , which `grid` converts to a range of vertical pixels on the graph.
  - ▶ legends and axes provide the inverse mapping


p2



## Components of a ggplot: coordinate system

- ▶ The coordinate system is another layer in how the data get mapped to the graphics device.
  - ▶ Usually Cartesian, but could be, e.g., polar coordinates, or a transformation.

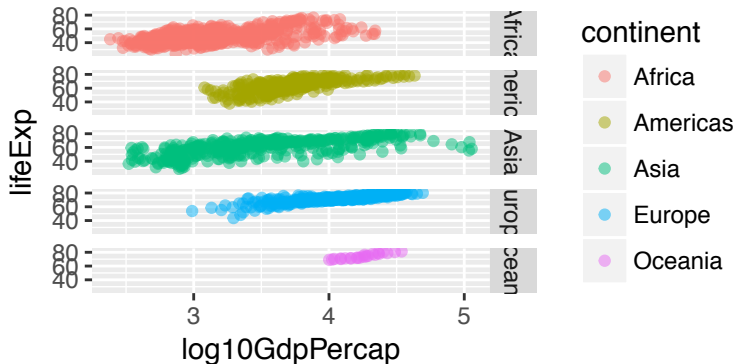
```
p2 + coord_trans(x="log10")
```



# Components of a ggplot: faceting

- How to break up the data into subsets and arrange multiple plots on the graphics device.

```
p2 + facet_grid(continent ~ .)
```



# Why so many components?

- ▶ A framework for the components of a graph.
- ▶ Gives the user the ability to change individual components one at a time.

# More details

- ▶ Layers
  - ▶ data, aesthetic mapping, geom, statistical transformation and position adjustment (to be defined)
- ▶ Tools for working with layers
- ▶ Scales, axes and legends
- ▶ Positioning: faceting and coordinate systems

## Example dataset: Diamonds

- ▶ Price and quality of about 54,000 diamonds.
- ▶ Quality measures are carat (size), cut, colour and clarity
- ▶ Also included are three measures of the dimensions of each diamond labelled x, y and z.

```
data(diamonds)
head(diamonds)
```

```
## # A tibble: 6 × 10
##   carat      cut color clarity depth table price      x      y      z
##   <dbl>    <ord> <ord>   <ord> <dbl> <dbl> <int> <dbl> <dbl> <dbl>
## 1  0.23    Ideal   E      SI2   61.5   55    326  3.95  3.98  2.43
## 2  0.21    Premium  E      SI1   59.8   61    326  3.89  3.84  2.31
## 3  0.23     Good   E      VS1   56.9   65    327  4.05  4.07  2.31
## 4  0.29    Premium  I      VS2   62.4   58    334  4.20  4.23  2.63
## 5  0.31     Good   J      SI2   63.3   58    335  4.34  4.35  2.75
## 6  0.24 Very Good  J     VVS2   62.8   57    336  3.94  3.96  2.48
```

# Layers

- ▶ We first initialize the plot and then start adding layers.
- ▶ Initializing is done with `ggplot()`.
  - ▶ We usually specify the default data and aesthetic mappings for all subsequent layers, though this is not necessary.
  - ▶ Without layers, the plot is not displayed.


```
p <- ggplot(diamonds, aes(x=carat, y=price, colour=cut))
```



## Adding layers

- ▶ Add with a +
- ▶ The `layer()` function can be used to specify a geom, stat and position
  - ▶ data and mapping will be inherited from initialization

```
p + layer(geom="point", stat="identity", position="identity")
```



# Shortcuts

- ▶ Shortcut functions are of the form `geom_XXX()` and `stat_XXX()`.
  - ▶ The `geom_XXX()` functions have a default stat and position
  - ▶ The `stat_XXX()` functions have a default geom and position
  - ▶ The `geom_XXX()` can over-ride the default stat and the `stat_XXX()` can over-ride the default geom though
- ▶ Call on the previous slide is equivalent to

```
p <- p + geom_point()
```

# Plot objects

- Notice that plot objects can be stored as R objects:

```
summary(p)
```

```
## data: carat, cut, color, clarity, depth, table, price, x, y, z
## [53940x10]
## mapping: x = carat, y = price, colour = cut
## faceting: <ggproto object: Class FacetNull, Facet>
##   compute_layout: function
##   draw_back: function
##   draw_front: function
##   draw_labels: function
##   draw_panels: function
##   finish_data: function
##   init_scales: function
##   map: function
##   map_data: function
##   params: list
##   render_back: function
##   render_front: function
##   render_panels: function
##   setup_data: function
##   setup_params: function
##   shrink: TRUE
##   train: function
##   train_positions: function
##   train_scales: function
##   vars: function
##   super: <ggproto object: Class FacetNull, Facet>
## -----
## geom_point: na.rm = FALSE
## stat_identity: na.rm = FALSE
## position_identity:
```